

CS 513
Solutions to #3

1. Suppose that you are given a *k-sorted* array, in which no element is farther than *k* positions away from its final (sorted) position. Give an algorithm which will sort such an array. Prove its correctness. Analyse its running time. Note: your algorithm should run faster than $\Theta(n \log n)$, that is, it should take advantage of the fact that the array is *k-sorted*.

Answer: The algorithm is as follows:

k-sort (*A*)

Divide input array *A* into blocks of size *k*: $B_i = A[(i-1)k+1] \cdots A[ik]$.

for $i = 1$ to n/k do

Mergesort(B_i, B_{i+1})

end

end algorithm

Running time: Each call to **mergesort** is on an array of size $2k$, so it takes $\Theta(k \log k)$. Since n/k such calls are made, the total time is $\Theta(n \log k)$.

Correctness:

Claim 1 After **mergesort**(B_i, B_{i+1}) has excuted, blocks B_1, \dots, B_i are correctly sorted and blocks $B_{i+1}, \dots, B_{n/k}$ are *k-sorted*.

Proof of claim: Base case: For $i = 0$, that is before we call **mergesort** at all, the claim is trivially true.

Inductive hypothesis: Suppose the claim is true for $i' < i$. By induction, B_1, \dots, B_{i-1} are correctly sorted. After calling **mergesort**(B_i, B_{i+1}), $B_{i+2}, \dots, B_{n/k}$ are still *k-sorted*. So we need to show that B_i is sorted and B_{i+1} is still *k-sorted*.

B_i is sorted because the elements that end up in B_i are the k smallest elements of B_i and B_{i+1} . **mergesort**(B_i, B_{i+1}) will corectly put the k smallest elements in their place.

Now, suppose that B_{i+1} is no longer *k-sorted*. Let j be the index of the rightmost element of B_{i+1} that is more than k away from its correct position. First, j cannot be the last position, since each element in B_{i+1} must end up in either B_{i+1} or B_{i+2} (and all these positions are within k of the last position). Now if there are l things greater than $A[j]$ in B_{i+1} , then $A[j]$'s final position cannot be within l of the right-hand side

of B_{i+2} . But then its final position is no more than k away, so $A[j]$ is within k of its final position, so B_{i+1} is k -sorted. \square

We conclude by noting that after calling **mergesort**($B_{n/k-1}, B_{n/k}$). We have that $B_1, \dots, B_{n/k-1}$ is sorted and $B_{n/k}$ is k -sorted. But clearly $B_{n/k}$ is also sorted, giving that the entire array is sorted. \square

2. Consider once again the k -sorted array of the previous problem. Show that any comparison based algorithm for sorting an almost sorted array makes $\Omega(n \log k)$ comparisons.

Proof: We must count the number of distinct k -sorted arrays. We will give a lower bound for this number. We will count all the k -sorted arrays so that elements need not cross blocks (as defined in algorithm above). We call all such arrays *block k -sorted*. To recap, in these arrays, when going from a block k -sorted array to a sorted array, the correct position for all elements in block B_i is still within block B_i .

Now, within each block, any ordering of the elements is possible, so there are $k!$ ways to order each block. That means that the number of distinct block k -sorted arrays is $(k!)^{n/k}$. So the number of k -sorted arrays is at least $(k!)^{n/k}$. So a decision tree for this problem has depth $\Omega(\log(k!)^{n/k}) = \Omega((n/k) \log(k!)) = \Omega((n/k)k \log k) = \Omega(n \log k)$. \square

3. Prove that if an n node graph has two of the following properties, it has the third:

- (a) It has $n - 1$ edges.
- (b) It is connected.
- (c) It is acyclic.

Answer: First show that (2),(3) imply (1).

Proof: Of course, by induction on n .

Base case: If $n = 1$, then trivially there are no edges, and $0 = 1 - 1$.

Inductive hypothesis: Suppose that each tree T' with $n' < n$ nodes has $n' - 1$ edges. Let T be an n node tree with e edges. Pick an arbitrary edge \overline{xy} in T . Then removing \overline{xy} leaves T disconnected into two trees. In particular, we know that x and y can no longer be connected, since if they were, adding \overline{xy} back would form a cycle in T . Any node v must use \overline{xy} in its path in T to either x or y (otherwise, we once again get a cycle), so we get two trees T_x and T_y , both of which have fewer than n nodes. T_x is the tree of nodes still connected x and T_y is the tree of nodes still connected to y . Let T_x have n_x nodes and e_x edges. Let T_y have n_y nodes and e_y edges. Then $e_x = n_x - 1$ and $e_y = n_y - 1$, by induction. Finally $n = n_x + n_y$ and $e = e_x + e_y + 1$ (the last one is for \overline{xy}), so $e = e_x + e_y + 1 = n_x - 1 + n_y - 1 + 1 = n - 1$, as desired. \square

Now show that (1),(3) imply (2).

Proof: Suppose otherwise. Let T be an acyclic graph with $n - 1$ edges and k connected components T_1, \dots, T_k . Let T_i have n_i nodes. Each T_i is connected and acyclic so it is a tree (from first part of proof). So it has $n_i - 1$ edges. The total edges for T is $\sum_{i=1}^k n_i - k = n - k$. Therefore $k = 1$ and the graph is connected. \square

Finally, (1),(2) imply (3).

Proof: Suppose otherwise. Let T be a connected graph with $n - 1$ edges, and with a cycle. Pick any edge from a cycle in T and remove it. The new graph T' is still connected, but it has $n - 2$ edges. If it still has a cycle, remove another edge from that cycle, and so on, until the remaining graph no longer has a cycle. This will be a tree, since it is still connected, but it will have at most $n - 2$ edges, thus contradicting the first part of this proof. \square

4. (a) Prove that the number of subtrees of a complete binary tree is not polynomial in the number of nodes.

Proof: We will underestimate the number of subtrees. Let T be an n leaf binary tree. It then has $2n - 1$ total nodes. Consider the class of subtrees S of T which has all the nodes in all levels but the bottom. In the bottom level, there is one subtree in S for each subset of the leaves. Clearly, each graph in S is connected, so it is a subtree of T , and clearly, T has at least $|S|$ subtrees. So the question is, how big is S ? But this is the same questions as asking, how many subsets of n leaves are there? There are 2^n such subsets, so $|S| = 2^n$. So the question now is, is there a constant c such that 2^n is $O(n^c)$ – this is what it means to be polynomial in n .

Claim 2 For any constant c , 2^n is not $O(n^c)$.

Proof of claim: $\lim_{n \rightarrow \infty} \frac{2^n}{n^c} = \lim_{n \rightarrow \infty} \frac{2^n}{2^{c \log n}} = \lim_{n \rightarrow \infty} 2^{n - c \log n} = \infty$.

- (b) Give an example of a class of trees $\{T_n\}$ where the number of subtrees is a polynomial in the number of nodes.

Answer: Let T_n be the class of n node chains. That is, if the nodes are $1, \dots, n$, then i is connected to $i - 1$ and $i + 1$, for $1 < i < n$. Clearly, T_n is connected and acyclic, so it's a tree. Now, any subtree must start at some i and end at some j and include all the nodes between, so there are $\binom{n}{2}$ subtrees for T_n . This is $\Theta(n^2)$.

5. Let F_i be the i th Fibonacci number (that is $F(0) = 0, F(1) = 1, F(n + 2) = F(n) + F(n + 1)$). Show that $F_{n+2} = 1 + \sum_{i=0}^n F_i$.

Proof: By induction, the base case $F_2 = 1$ is trivial.

The inductive hypothesis is that the claim holds for all $n' < n + 2$. Then $F_{n+2} = F_{n+1} + F_n = F_n + 1 + \sum_{i=0}^{n-1} F_i = 1 + \sum_{i=0}^n F_i$. \square

6. Show that if you have a polynomial time algorithm for Hamiltonian Path, that you have a polynomial time algorithm for sorting.

Answer: We already have a polynomial time algorithm for sorting, e.g. $O(n \log n)$ for mergesort. Thus, the implication we must prove is vacuously true. \square

7. The *Bounded Degree Spanning Tree* (BDST) problem is the following:

Input: Graph G and integer k .

Output: Yes, if G has a spanning tree where every node has degree at most k , No, otherwise.

Suppose there is no polynomial time algorithm for Hamiltonian Path. Show that there is no polynomial time algorithm for BDST.

Proof: We must give a poly time algorithm for HamP using BDST. Suppose we want a Hamiltonian path from x to y in G . Create G' from G by adding nodes x' and y' and hooking x' to x and y' to y . Now $HamP(G, x, y)$ is true iff $BDST(G', 2)$ is true. Why, because a spanning tree of degree two is exactly a hamiltonian path, and x' and y' for the spanning tree to start and end at these nodes. \square